

ECE437/CS481

M03A: CPU SCHEDULING
SCHEDULING CONCEPT & ALGORITHMS

CHAPTER 6.1-6.3

Xiang Sun

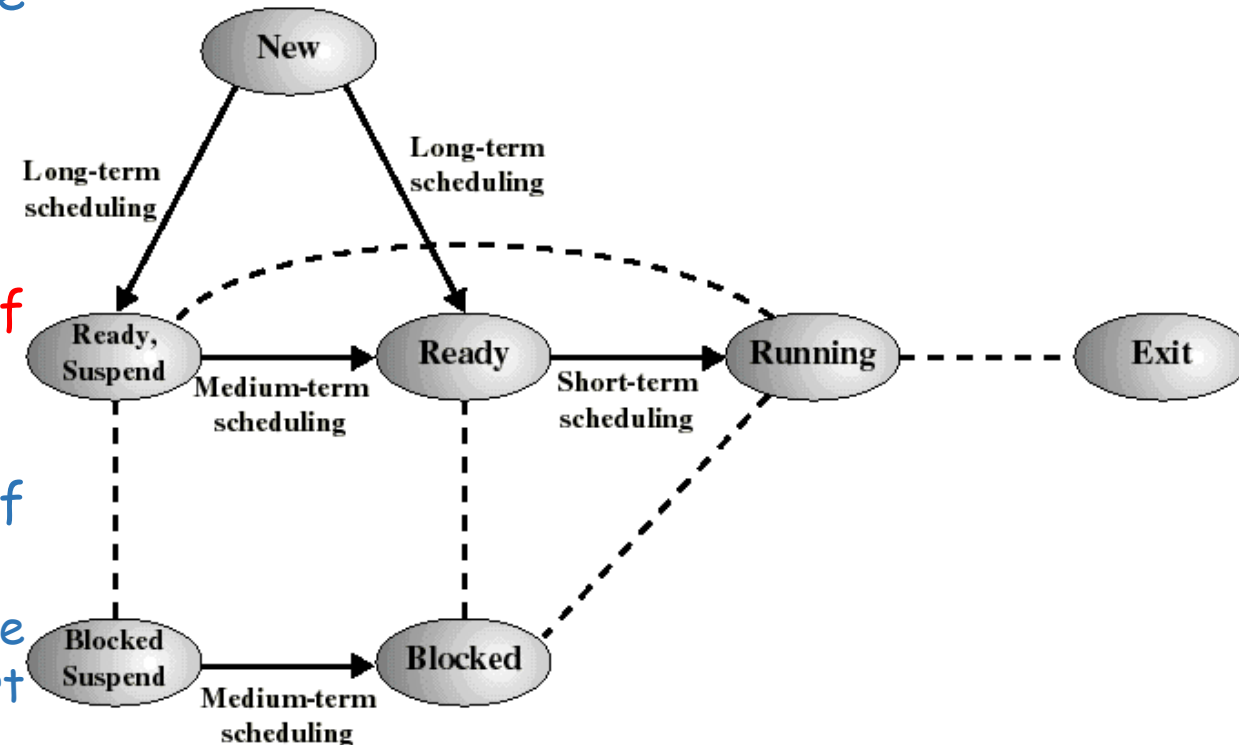
The University of New Mexico

A decorative blue wavy line that spans the width of the slide, starting with a small upward curve on the left, dipping into a V-shape in the center, and then curving back up on the right before continuing as a straight line to the edge.

Levels of Scheduling

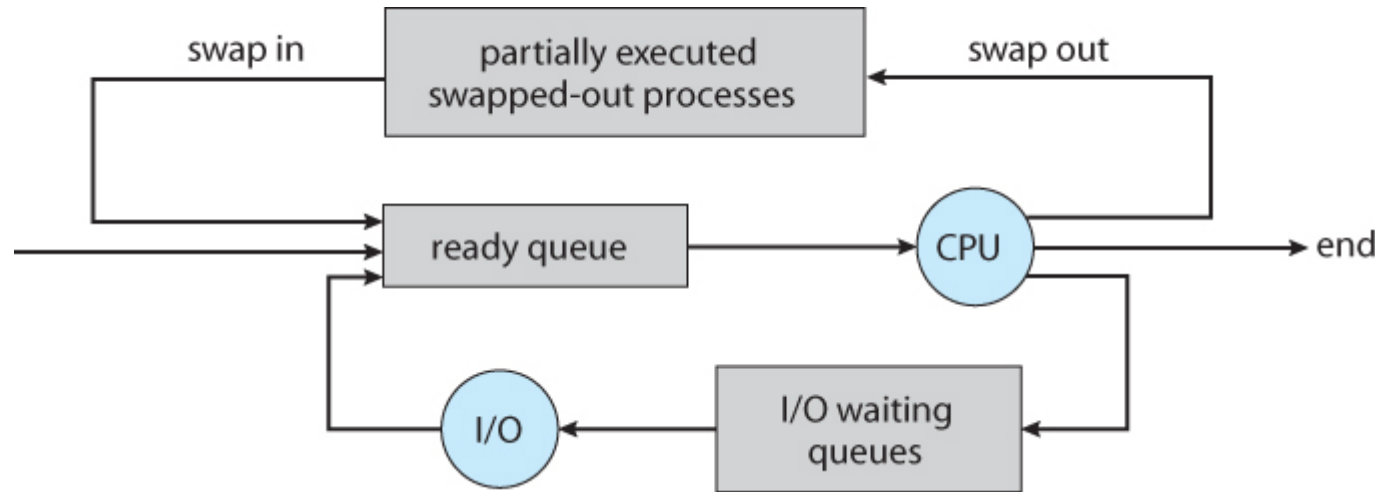
□ Long-term scheduling/Job scheduler (High-level)

- Determine which process should be brought into the ready queue
- Invoked infrequently
- Coarse-grained control of the **degree of multiprogramming**
- Should balance different types of processes:
 - **I/O-intensive process**: spends more time doing I/O than computations, many short CPU bursts.
 - **CPU-intensive process**: spends more time doing computations; few very long CPU bursts.



Levels of Scheduling

❑ Medium-term scheduling (swapping in/out)

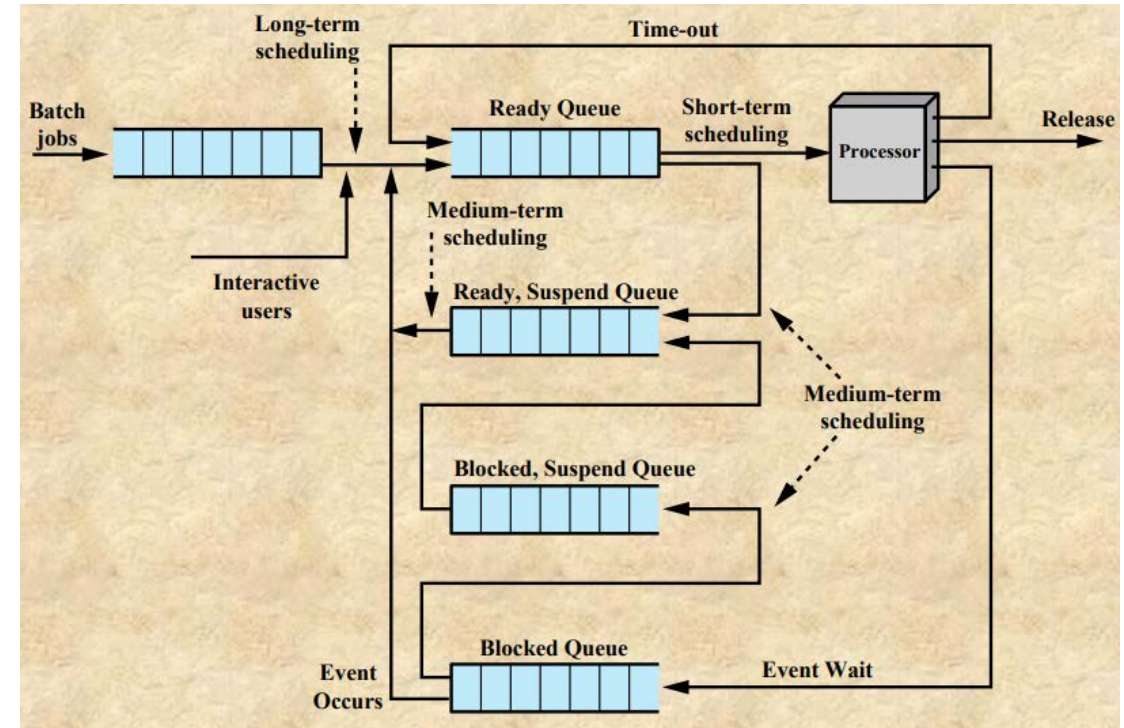


- Adjust the degree of multiprogramming by **swapping**.
- Swapping: removes a process from memory, stores on disk, and bring back in from disk to continue execution later on.
 - ✓ **Swap out**: memory-to-disk
 - ✓ **Swap in**: disk-to-memory
- Normally, the medium-term scheduler may decide to swap out a process which has not been active for some time, or a process which has a low priority, or a process which is taking up a large amount of memory in order to free up main memory for other processes, etc.

Levels of Scheduling

❑ Short-term scheduling/CPU scheduling

- Determine which process in the ready queue will be executed by the CPU
- Invoked when an event occurs that may lead to the blocking of the current process or that may provide an opportunity to preempt a currently running process in favor of another
 - ✓ Clock interrupts
 - ✓ I/O interrupts
 - ✓ Operating system calls
 - ✓ Signals, etc.
- Scheduler is consuming CPU too! It executes most frequently, must be very careful about its computation overhead.



Scheduling Objectives

□ Metrics—system-wide

- Maximize processor/CPU utilization
 - ✓ percentage of time that CPU is running users' processes, to keep system as busy as possible
- Maximize throughput
 - ✓ number of processes completed per time unit
 - ✓ number of instruction executed per time unit
- Fairness
 - ✓ don't starve any processes—treat the all the same



Scheduling Objectives

□ Metrics—per process/user-oriented

- Minimize waiting time
 - ✓ time spent in **ready queue**
- Minimize turnaround time
 - ✓ for a batch job, time between the submission of a process and the completion of the process, equal to sum of waiting time in ready queue and blocking queue, plus running time.
- Minimize response time
 - ✓ for interactive job, time from the submission of a request until **the first response happens**.



Scheduling Objectives

□ Metrics

- Achieve a balance between response and utilization
- Minimize overhead (system level)
 - ✓ Context switching
 - ✓ Scheduling complexity
- It is difficult to find the optimal solution of the scheduling since the scheduling problem is mostly an NP-hard/NP-complete problem (e.g., job shop scheduling). Instead, looking for heuristic approaches is the common way.



❑ Non-preemption V.S. Preemption

➤ Non-preemption

- ✓ once a process is scheduled to CPU, it continues execution until either it terminates, or relinquish (e.g. yield), or switches to a blocked state.

➤ Preemption

- ✓ OS allowed to reallocate CPU (e.g. timeout, for higher priority process, etc.) to other process.

❑ Concurrency V.S. Parallelism

➤ Concurrency

- ✓ schedule multiple processes onto a single CPU—time multiplex manner

➤ Parallelism

- schedule multiple processes onto multiple CPUs—spatial multiplex

Process Scheduling

Scheduling algorithm

Policy: to decide who gets to run

+

Dispatcher

Mechanism: how to do the context switch

- ❑ Decision mode--when to do the scheduling
 - non-preemption
 - preemption—high **priority** or periodically (the process's time slice expires)
- ❑ Priority function
 - static information, e.g., job length, CPU-intensive or I/O intensive, memory requirement
 - dynamic information, e.g., deadline, arrival time, recent CPU consumption

Process Scheduling

Scheduling algorithm

Policy: to decide who gets to run

+

Dispatcher

Mechanism: how to do the context switch

□ Events affect the scheduling algorithm:

- 1) Current process goes from **running to waiting state** (e.g., wait for I/O)
- 2) Current process terminates, **running to exit state**
- 3) Current process goes from **running to read state** (e.g., time slice is up)
- 4) I/O is complete --- Current process goes from **waiting to ready**

Process Scheduling

Scheduling algorithm

Policy: to decide who gets to run

+

Dispatcher

Mechanism: how to do the context switch


- ❑ Dispatcher module gives control of the CPU to the process selected by the **short-term scheduler**; this involves:
 - 1) switching context
 - 2) switching between user mode and kernel mode
 - 3) jumping to the proper location in the user program to restart that program
- ❑ **Dispatch latency**: time consumption of the dispatcher to stop one process and start another.

Scheduling Algorithms

❑ First-Come, First-Served (FCFS)

- General specification
 - ✓ Decision mode: non-preemption Priority: equal
- As a process become ready, it join the ready queue, scheduler always selects process from the front of the ready queue.

Process	Arrival time	Service time (Burst time)
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



Waiting time for P1 = 0; P2 = 1; P3 = 5, P4 = 7, P5 = 10,
Average waiting time: $(0 + 1 + 5 + 7 + 10)/5 = 4.6$

Scheduling Algorithms

❑ First-Come, First-Served (FCFS)

- Simple to implement; low overhead, since no priority calculation, no extra context switch.
- Average waiting time may be long and suffer from **convoy effect**.
 - ✓ Convoy effect: long waiting time due to the slow processes.

Case 1:

- ✓ suppose that the processes arrive in the order :P₁, P₂, P₃.

- ✓ The Gantt Chart for the schedule is



- ✓ Waiting time P₁=0; P₂=24; P₃=27
- ✓ Average waiting time: $(0+24+27)/3=17$

Process	Service time
P ₁	24
P ₂	3
P ₃	3

Case 2:

- ✓ suppose that the processes arrive in the order :P₂, P₃, P₁.

- ✓ The Gantt Chart for the schedule is



- ✓ Waiting time P₁=6; P₂=0; P₃=3
- ✓ Average waiting time: $(6+0+3)/3=3$

Scheduling Algorithms

❑ Shortest Job Next (SJN)/Shortest Job First(SJF)

➤ General specification

- ✓ Decision mode: non-preemption by default (could be implemented as preemption)
- ✓ Priority: $1/\text{service time}$ —a process with a lower service time has a higher priority

➤ Pros:

- ✓ Decrease average waiting time as compared to FCFS

➤ Cons:

- ✓ Higher complexity as compared to FCFS
- ✓ Difficult to predict service time
- ✓ Risk to **starving** longer process, as long as there are shorter processes around

❑ Shortest Remaining-time First (SRF)

➤ General specification

- ✓ Decision mode: **preemption**
- ✓ Priority: $1/(\text{service time} - \text{running time})$


➤ Still need estimate of service time, and can starve longer processes.

➤ May give better turnaround time than SJN/SJF.

Scheduling Algorithms

❑ Shortest Job Next (SJN)/Shortest Job First(SJF)

- Suppose that the processes arrive in the order: P1 , P2 , P3 , P4
- The Gantt Chart for the schedule is:




Process	Service time
P1	6
P2	8
P3	7
P4	3

- Waiting time for P1=3 , P2=16 , P3=9 , P4=0
- Average waiting time: $(3+16+9+0)/4=7$
- Suppose FCFS is applied, average waiting time: $(0+6+14+21)/4=10.25$

Scheduling Algorithms

Shortest Job Next (SJN)/Shortest Job First(SJF)



Process	Arrival time	Service time
P1	0	3
P2	2	6
P3	4	4
P4	6	5
P5	8	2

- Waiting time for P1=0 , P2=1, P3=7 , P4=9, P5=1.
- Average waiting time: $(0+1+7+9+1)/5=3.6$

❑ Shortest Job Next (SJN)/Shortest Job First(SJF)

- How to predict the service/burst time of a process?
- In reality, the prediction is done by using the length of previous service/burst time of the process.
 1. t_n = actual length of n^{th} CPU burst
 2. τ_{n+1} = predicted value for the next CPU burst
 3. $\alpha, 0 \leq \alpha \leq 1$
 4. Define : $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$.

□ Shortest Job Next (SJN)/Shortest Job First(SJF)

$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$$

➤ If $\alpha=0$

- ✓ $\tau_{n+1} = \tau_n$
- ✓ No weightage to recent history

➤ If $\alpha=1$

- ✓ $\tau_{n+1} = \alpha t_n$
- ✓ Only the actual last CPU burst counts

✓ If $0 < \alpha < 1$, we expand the formula:

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \alpha t_{n-1} + \dots + (1 - \alpha)^j \alpha t_{n-j} + \dots + (1 - \alpha)^{n+1} \tau_0$$


- ✓ Since both α and $(1 - \alpha)$ are less than or equal to 1, each successive term has less weight than its predecessor

Scheduling Algorithms

❑ Shortest Remaining-time First (SRF)

➤ Suppose we have four processes arriving in the order showing in the table.

➤ The Gantt Chart for the schedule is:




Process	Arrival time	Service time
P1	0	6
P2	1	2
P3	2	8
P4	3	4

➤ Waiting time for P₁=6 , P₂=0 , P₃=10 , P₄=0

➤ Average waiting time: $(6+0+10+0)/4=4$

➤ Suppose SJN is applied, average waiting time: $(0+5+10+5)/4=5$



Scheduling Algorithms

□ Round Robin (RR)


- General specification
 - ✓ decision: preemption, periodically with time slice
 - ✓ priority: equal
- Use preemption based on clock - time slicing (time quantum), generate interrupt at periodic intervals, usually 10-100ms
- When interrupt occurs, place running process back to Ready queue, select next process to run by using **FCFS**.
- Designed especially for **interactive jobs**.
- Average waiting time may be long.
- What's the right length of a time slice?
 - ✓ short means short processes move through quickly, but high overhead to deal with process scheduling and context switching.
 - ✓ very long time slice making it degenerate into FCFS.
 - ✓ should be slightly greater than time of typical job burst time.

Scheduling Algorithms

□ Round Robin (RR)

- Suppose that the processes arrive in the order: P1 , P2 , P3. The time quantum length is 4.

- The Gantt Chart for the schedule is:




Process	Service time
P1	24
P2	3
P3	3

- Wait time for P1=6; P2=4; P3=7
- Average waiting time: $(6+4+7)/3=5.67$.
 - ✓ Recall that the average waiting time for FCFS is 17. However, Ave waiting time (RR) < Ave waiting time (FCFS) is not always true.
- Typically, RR incurs higher average turnaround time, but lower response time than FCFS.

Scheduling Algorithms

Round Robin (RR)



Waiting time for P1 = 1; P2 = 10; P3 = 10, P4= 9, P5 = 5,

Average waiting time: $(1 + 10 + 10 + 9 + 5)/5 = 7$

Recall that the average waiting time for FCFS is 4.6

Process	Arrival time	Service time (Burst time)
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2

□ Round Robin (RR)

➤ Assumption of RR

- ✓ there are n processes in the ready queue, and
- ✓ the time quantum length is q

➤ Then

- ✓ each process gets $1/n$ of the CPU time in chunks of **at most** q time units at once.
- ✓ No process waits more than $(n-1)q$ time units.


➤ Performance

- ✓ If q large ($q \rightarrow +\infty$), $RR \rightarrow FIFS$
- ✓ If q is small, RR incurs frequent process context switch, thus generating high overhead.

Scheduling Algorithms

□ Round Robin (RR)

➤ Time Quantum and Context Switch



Scheduling Algorithms

Round Robin (RR)

- Time quantum length selection.
 - ✓ Must be substantially larger than the time required to handle the clock interrupt and dispatching (i.e., dispatch latency).
 - ✓ Should be larger than the typical length of interaction (but not much more) to avoid penalizing I/O intensive processes.

