

ECE437/CS481

M03A: CPU SCHEDULING
MORE ABOUT SCHEDULING

CHAPTER 6.4-6.8

Xiang Sun

The University of New Mexico

A decorative blue wavy line that spans the width of the slide, starting with a small upward curve on the left, dipping into a V-shape in the center, and then curving back up on the right before continuing as a straight line to the edge.

Dynamics with feedback

❑ Feedback in CPU scheduling

- Service/burst time of a process may vary over time.
- Processes may be terminated and create by the system over time.
- Feedback is used to dynamic adjust the priorities of different processes.



Dynamics with feedback

- Feedback to emphasize on **fairness**
 - For N users
 - ✓ Promise every user about $1/N$ of CPU time
 - Dynamically bookkeeping two variables:
 - ✓ Promised time: T_p
 - ✓ Actual time used: T_u
 - Let priority function $F = T_u/T_p$
 - ✓ The lower the value, the higher in priority
 - ✓ case $F = 1$, just right, kept promise
 - ✓ case $F < 1$, under provisioned, the processes can request more CPU time.
 - ✓ case $F > 1$, over provisioned, the processes have to slow down.

Dynamics with feedback

□ Feedback to emphasize on **fairness**

➤ Fair enough?

- ✓ For those didn't get a fair share of CPU time → priority raised
- ✓ If new I/O-intensive processes (which require low CPU time/service time) regularly enter the system, computation-intensive processes potentially get starved, increasing their turnaround time.
- ✓ Achieving fairness is difficult.



Dynamics with feedback

□ Feedback to emphasize on **aging**

- Favor a user who spent more waiting time at Ready queue.
- Dynamically bookkeeping two variables:
 - ✓ Total waiting time: T_w
 - ✓ Actual time used: T_u
- Let priority function $F = T_u/T_w$
 - ✓ The lower the value, the higher in priority



Classification of processes

□ By the nature of requirement

- **Foreground processes**—the system/shell may wait for the processes to finish
 - ✓ I/O-intensive processes
- **Background processes**-- the system/shell does not have to wait for the processes to finish before it can run more processes
 - ✓ computation-intensive processes

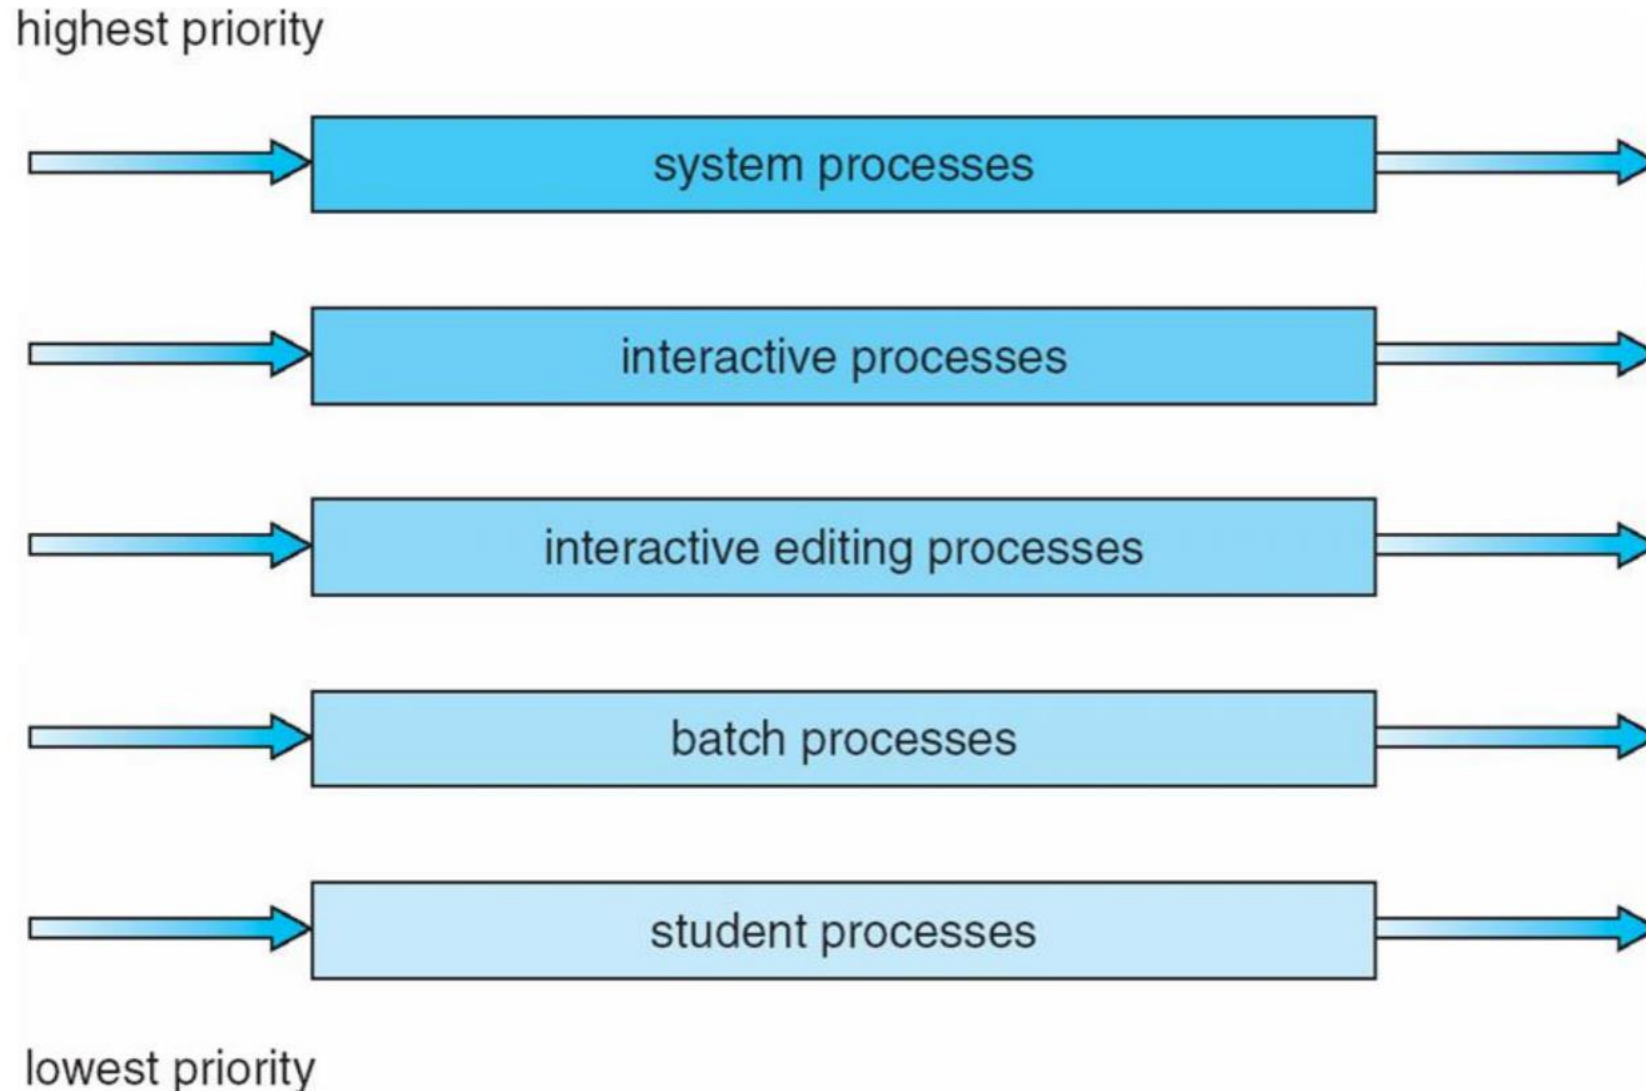
□ By run-time characteristics

- **System process**
 - ✓ Accomplishing system task: e.g., paging
 - ✓ A user process executing kernel task via system call, e.g., I/O request
- **User processes**

Multilevel Queues

- ❑ Ready queue is partitioned into **separate queues**
- ❑ Each queue has its **own scheduling algorithm**
 - Typical example
 - ✓ foreground queue- RR
 - ✓ background queue- FCFS
- ❑ **Scheduling** must be done among these queues:
 - Fixed priority scheduling; (i.e., serve all from foreground, then from background). Possibility of starvation.
 - Time slice - each queue gets a certain amount of CPU time, which it can schedule among its processes; e.g., 80% to foreground in RR, 20% to background in FCFS.

Multilevel Queues



Multilevel Queues

❑ Multilevel queues + feedback

- Process can move among queues based on their feedbacks (e.g., waiting time or burst time).

❑ Design & implementation of multilevel queues

- # of queues and their related scheduling strategies (i.e., intra-queue scheduling).
- Scheduling strategies among different queues (i.e., inter-queue scheduling).
- Method used to adjust the process among queues (feedback control).
- Initially process deployment (i.e., determine which queue a process will enter initially)

Multilevel Feedback Queues

□ An example of multilevel feedback queues

➤ There are three queues, and their **intra-queue scheduling** are:

- ✓ Q0 - RR with time quantum 8 ms
- ✓ Q1 - RR with time quantum 16 ms
- ✓ Q2 - FCFS


➤ The **inter-queue scheduling** is

- ✓ Fix priority scheduling: Q0—high, Q1—medium, Q2—low.

➤ The **feedback control** are designed as follows:

- ✓ A new process first enters queue Q0
- ✓ If the process does not finish in 8 ms, the process is moved to Q1; otherwise, stay in Q1.
- ✓ Once the process moves to Q1, it will receive 16 ms in the next cycle.
- ✓ If the process still does not finish in 16 ms, it is preempted and moved to Q2.

➤ I/O-intensive process will normally end up on high priority queue (Q0), and computational-intensive process will normally end up on low priority (Q2).

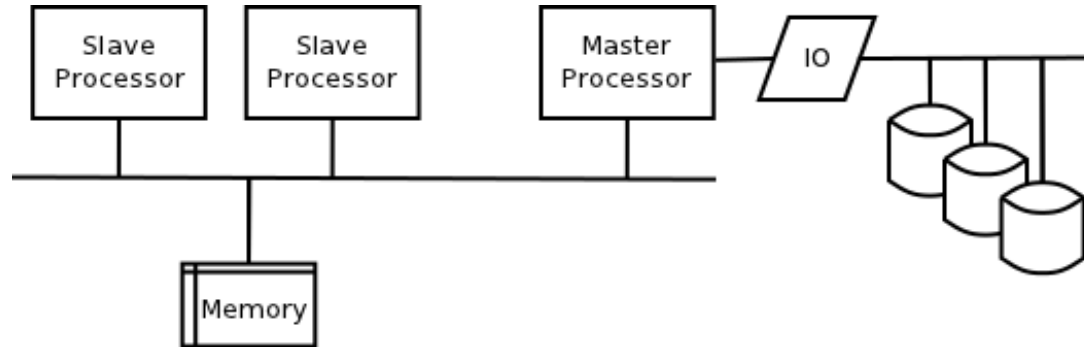


Multiprocessor Scheduling

- ❑ What is multiprocessor scheduling
 - Given a set of runnable processes/threads, and a set of CPUs, assign processes/threads to CPUs
- ❑ Same metrics as uniprocessor scheduling
 - Fairness, efficiency, throughput, response time...
- ❑ But also new considerations
 - Relationship among processors
 - Load balancing
 - **Processor affinity**—keep a process running on the same CPU

Asymmetric Multiprocessor Processing (AMP)

❑ Solution 1: Asymmetric multiprocessor processing (Centralized processing)

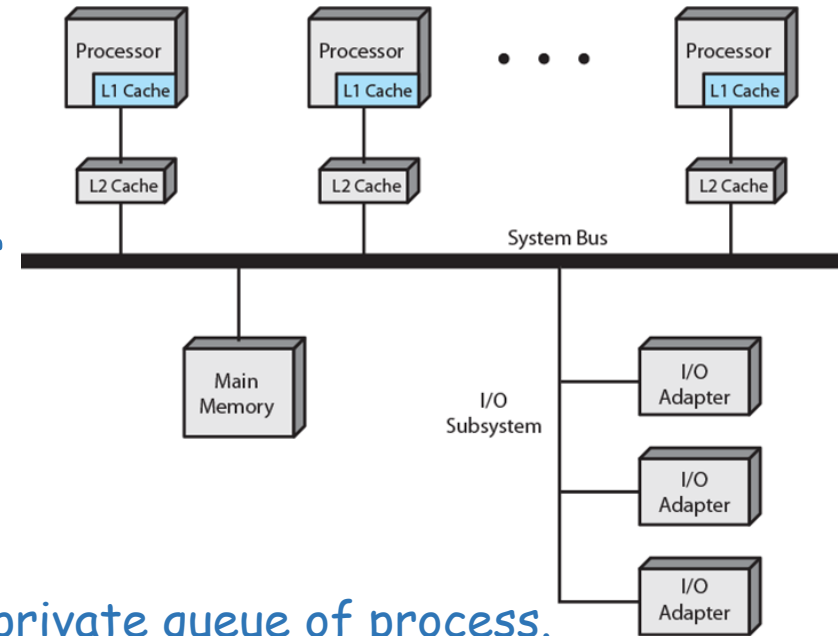


- Two types of processors (one master processors and a number of slave process). One ready queue.
- Master processor makes scheduling decision and handles I/O operations.
- Slave processors executes assigned processes. The master processor would balance the workload among slave processes
- If a master processor fails, a slave processor become the master processor. If a slave processor fails, its allocated processes are switched to other slave processors.

Symmetric Multiprocessor Processing (SMP)

❑ Solution 2: Symmetric multiprocessor processing (Distributed processing)

- Each processors is self-scheduling.
- All processes may be in a common ready queue (global ready queue) or each processor may have its own private queue for ready processes.
- Balancing the workload among processors is necessary to maximize the performance of the system.
 - ✓ Load balancing is necessary only if each processor has its own private queue of process.
 - ✓ Two general approaches to achieve load balancing: **Push migration** and **Pull migration**.
 - ✓ Push migration: A surveillance task periodically checks the workload on each processor and moves processes from processors with high load to processors with low load if needed.
 - ✓ Pull migration: A processor's scheduler notices its queue is empty (or shorter than a threshold), and tries to fetch a process from another processor.



Processor affinity in SMP

□ Recall

- Processors share main memory.
- Processors have their own local cache memories.
- Recently accessed data are stored in local cache memories in order to speed up data retrieval.

□ Process affinity

- Try to keep the process running in the same processor.
- Benefit of process affinity: quicker to restart process on same processor since the cache may already contain needed data.
- Two types of processor affinity:
 - ✓ Soft Affinity: a scheduler has a policy of attempting to keep a process running on the same processor but not guaranteeing it will do so.
 - ✓ Hard Affinity: some systems such as Linux have a **system call** to specify that a process shall execute on a specific processor.

```
unsigned long mask = 7; /* processors 0, 1, and 2 */
unsigned int len = sizeof(mask);
if (sched_setaffinity(0, len, &mask) < 0) {
    perror("sched_setaffinity");
}
```

Realtime System

- ❑ In a **realtime** system, the correctness of the system depends on
 - The logical result of the computation, AND.
 - The time when the results are produced.
- ❑ **Deadline** is associated a particular task
 - **hard deadline**: required to complete a critical task within a guaranteed amount of time
 - **soft deadline**: the deadline is desirable but not mandatory; still make sense to schedule and complete the task even if it has passed its deadline.
- ❑ Two types of processes in realtime system, i.e., Periodic and aperiodic process
 - Periodic processes: arriving at fixed frequency, can be characterized by 3 parameters (C,D,T) where C = service/burst time, D = deadline, T = period (e.g. 20ms, or 50HZ). Periodic processes are called Time-driven processes, their activations are generated by timers.
 - Aperiodic processes : all processes that are not periodic, also known as event-driven, their activations may be generated by external interrupts.

□ Deadline scheduling

- Need preemptive strategy & priority function
- The service time of each task is known in advance.
- Complete task neither too early nor too late.
- Find a schedule for all the tasks so that each meets its deadline.
 - ✓ What if we can't schedule all task?
 - ✓ What if the run-time is not constant but has a known probability distribution?
- A popular algorithm: EDF (**Earliest Deadline First**)


Earliest Deadline First (EDF) Scheduling

- ❑ Scheduler selects a job with EDF
 - ✓ The highest priority job is the one with the earliest **absolute deadline**;
 - ✓ If two jobs have the same deadline, chose one of the two at random;
 - ✓ The priority of a job is dynamic since the job's deadline changes over time;
 - ✓ Decision mode: preemption (by default).

Earliest Deadline First (EDF) Scheduling

□ Example: scheduling with EDF

- ✓ Three processes τ_1, τ_2, τ_3 are initially in the ready queue.
- ✓ $\tau_1=(1,4)$, which indicates that the burst time of τ_1 is 1 time unit, and τ_1 will be in the ready queue after each 4 time unit; the relative deadline of τ_1 is also 4 time unit. Accordingly, $\tau_2=(2,6)$, and $\tau_3=(3,8)$.



$$CPU\ utilization = \frac{1}{4} + \frac{2}{6} + \frac{3}{8} = \frac{23}{24}$$

Earliest Deadline First (EDF) Scheduling

□ **Theorem:** Given a job set of periodic or sporadic jobs, with relative deadlines equal to periods, the job set is schedulable by EDF iff

$$U = \sum_{i=1}^N \frac{C_i}{T_i} \leq 1$$

where C_i is the burst time of job i , T_i is the relative deadline of job i , N is the total number of jobs of the job set, and U is the CPU utilization.


➤ **Lemma:** EDF is an optimal algorithm, in the sense that if a job set is schedulable, then it is schedulable by EDF.

- ✓ If $U > 1$, no algorithm can successfully schedule the job set;
- ✓ If $U \leq 1$, EDF can always provide a feasible schedule.

Earliest Deadline First (EDF) Scheduling

□ Domino effect with EDF

- If $U > 1$ (i.e., the job is NOT schedulable), we have the domino effect with EDF: it means that all tasks miss their deadlines.
- An example of domino effect:
 - ✓ Four processes: τ_1 , τ_2 , τ_3 , and τ_4 .
 - ✓ $\tau_1 = (2, 5)$, $\tau_2 = (2, 6)$, $\tau_3 = (2, 7)$, $\tau_4 = (2, 8)$




Rate Monotonic(RM) Scheduling

□ Rate Monotonic (RM) Scheduling

- The highest priority job is the one with the earliest **relative deadline**/smallest period;
- If two jobs have the same deadline, chose one of the two at random;
- The priority of a job is static;
- Decision mode: preemption (by default).

- ✓ τ_1 and τ_2 never miss their deadlines;
- ✓ τ_3 misses a lot of deadline;
- ✓ τ_4 is not executed!



Rate Monotonic(RM) Scheduling

□ Rate Monotonic (RM) Scheduling

- RM cannot guarantee all the jobs in the job set meet their deadlines, even if the job set is schedulable (i.e., $U \leq 1$).
- $\tau_1=(1,4)$, $\tau_2=(2,6)$, and $\tau_3=(3,8)$. $U = \frac{1}{4} + \frac{2}{6} + \frac{3}{8} = \frac{23}{24}$

