

ECE437/CS481

M06C: VIRTUAL FILE SYSTEM & NETWORK FILE SYSTEM

CHAPTER 12.8

Xiang Sun

The University of New Mexico

A decorative blue wavy line that spans the width of the slide, starting with a thin line, dipping into a V-shape, and then rising back to a thin line, creating a stylized horizon or wave effect.

Virtual File System

□ Lots and lots of file system types

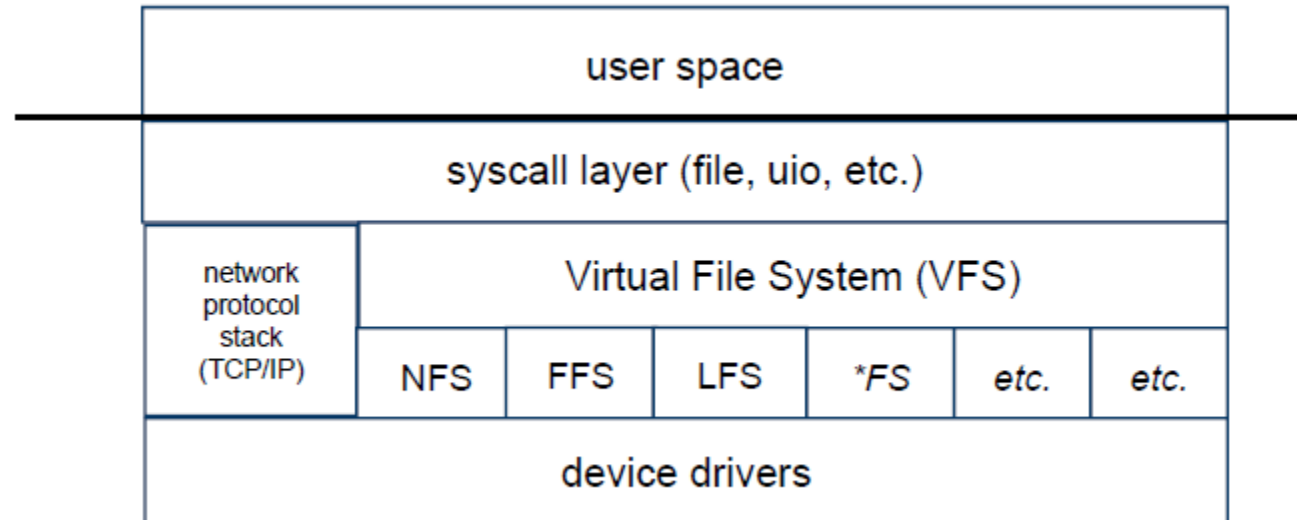
- Different file systems may have different directory structures, block sizes, etc.
- Is it possible to support more than one file system types on a single OS?

File System	Windows XP	Windows 7/8/10	macOS (10.6.4 and earlier)	macOS (10.6.5 and later)	Ubuntu Linux	Playstation 4	Xbox 360/One
NTFS	Yes	Yes	Read Only	Read Only	Yes	No	No/Yes
FAT32	Yes	Yes	Yes	Yes	Yes	Yes	Yes/Yes
exFAT	Yes	Yes	No	Yes	Yes (with ExFAT packages)	Yes (with MBR, not GUID)	No/Yes
HFS+	No	(read-only with Boot Camp)	Yes	Yes	Yes	No	Yes
APFS	No	No	No	Yes (macOS 10.13 or greater)	No	No	No
EXT 2, 3, 4	No	Yes (with third-party software)	No	No	Yes	No	Yes

Virtual File System

❑ Virtual file system

- A kernel software layer that handles all system calls related to file systems. Its main strength is providing a **common interface** to several kinds of file systems.



- ✓ Processes make system calls to VFS rather than specific file system interfaces.
- ✓ VFS separates file-system generic operations from implementation details, and then dispatches operations to appropriate file system implementation routines.

Virtual File System

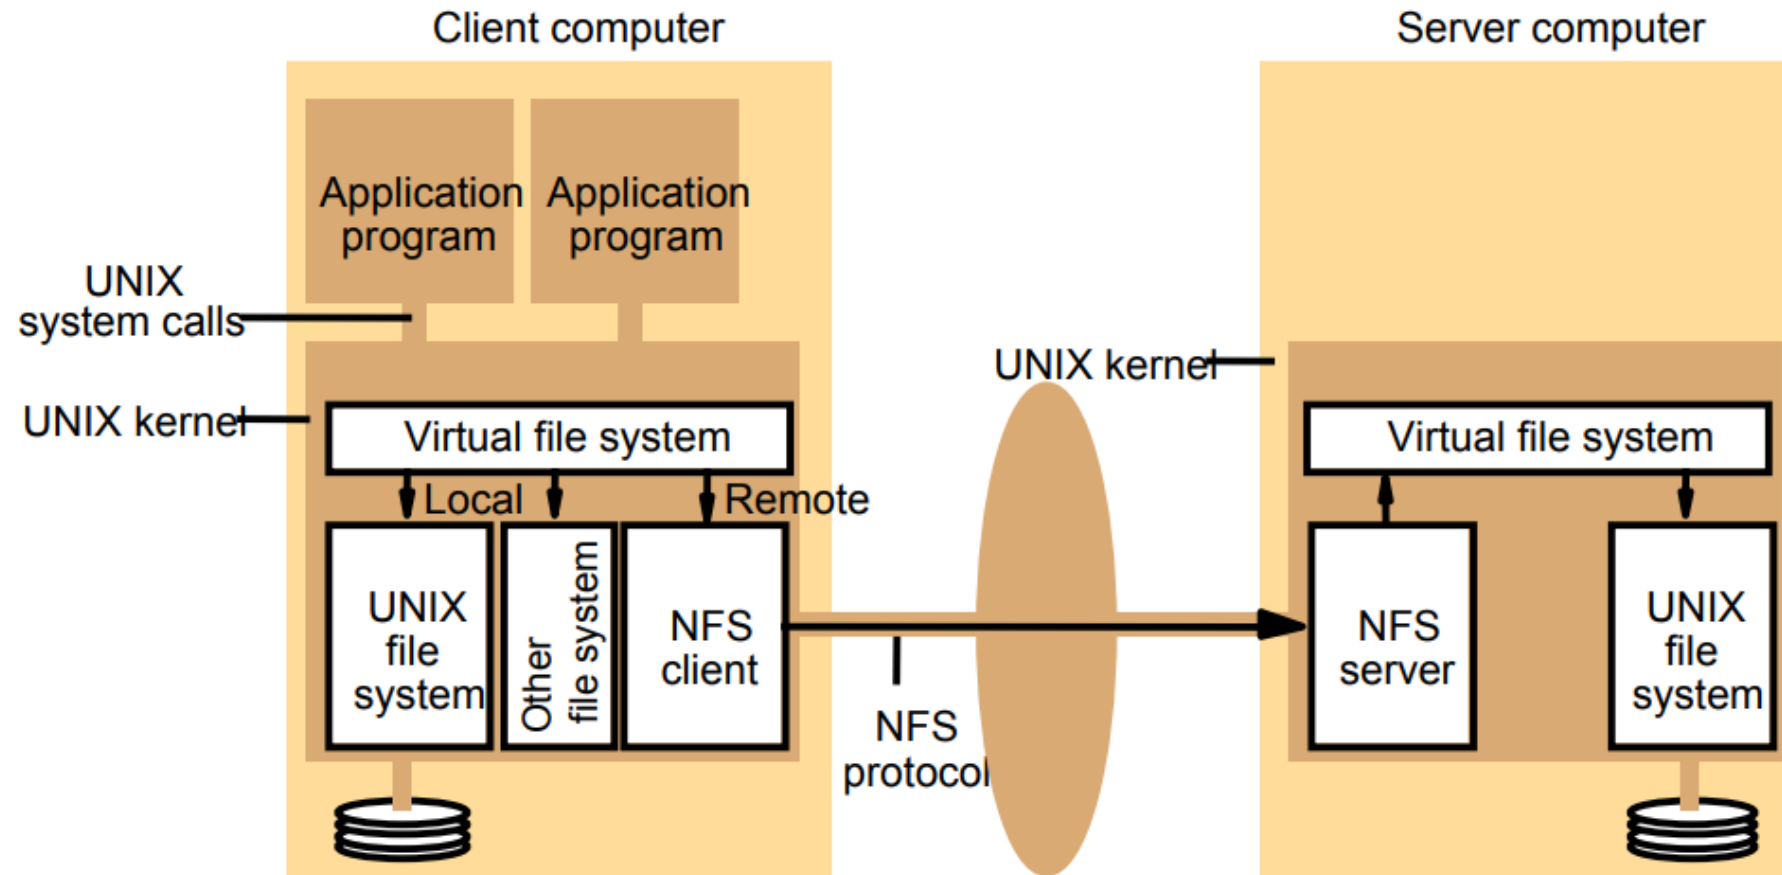
□ VFS system calls (~100)

- **filesystem** operations - mounting, info, chroot, pivot_root
- **directory** operations - chdir, getcwd, link, unlink, rename, symlink
- **file** operations - open/close, read/write, stat, permissions, seek
 - chmod, chown, stat, creat, umask, dup, fcntl, truncate
 - read/write, readv/writev, pread/pwrite
- **memory mapping** files - mmap, munmap, madvise, mlock
- **wait** for input - poll, select
- **flushing** - synch, fsync, msync, fdatasync
- **file locking** - flock

Network File System

- ❑ Purpose: access disks attached to other computers
- ❑ NFS is a protocol for remote access to a file system
 - Does not implement a file system per se
 - NFS is a network protocol layer over TCP/UDP
 - ✓ Original implementations use UDP as a transport layer protocol for low overhead.
 - ✓ Some newer implementations use TCP as a transport layer protocol for reliable communications.
 - Remote access is transparent to applications
- ❑ File system, OS, and architecture independent
- ❑ Implemented as a client/server architecture
 - Local file system (i.e. client) requests are forwarded to a remote server
 - Requests are implemented as remote procedure calls (RPCs)

Network File System





- ❑ NFS server check identity (e.g., uid, gid) for each access.
- ❑ NFS server is stateless, doesn't keep open files for clients
 - ✓ In NFS v4, stateful server are used

Network File System

□ Communications in NFS

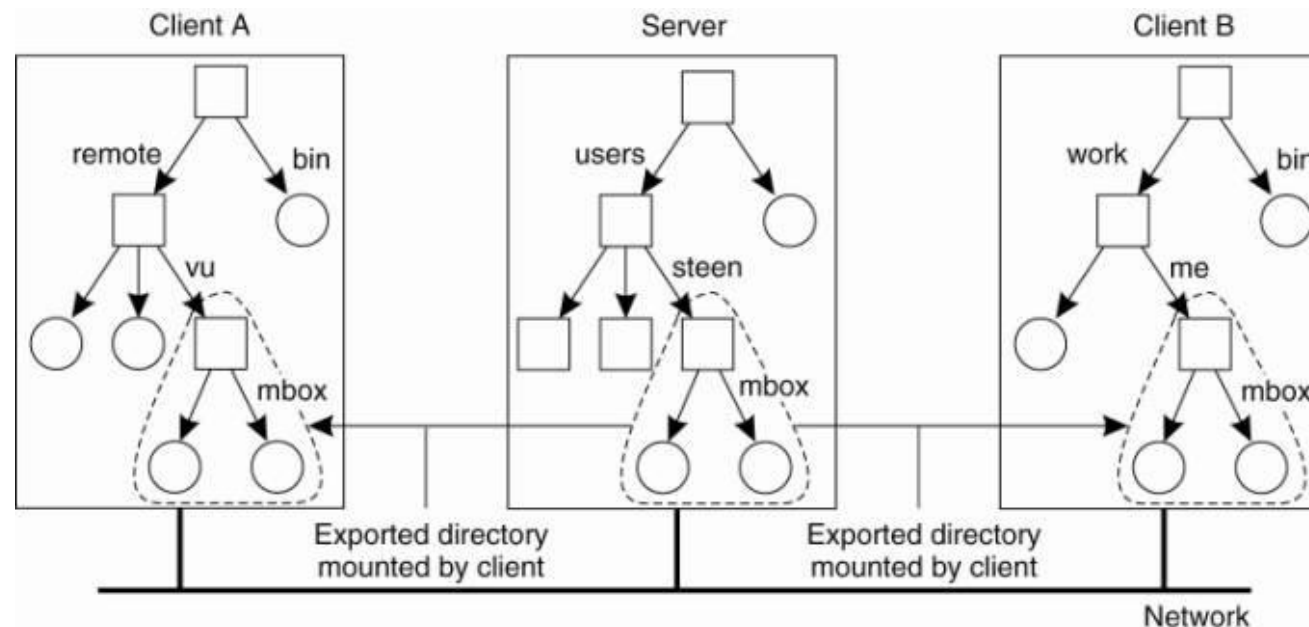
- ✓ The communications between a server and a client is based on the Open Network Computing RPC protocol.
- ✓ Every NFS operation (e.g., `open()`, `read()`, `write()`) can be implemented as a single remote procedure call to a file server.
- ✓ In order to reduce the communications overhead (i.e., reduce the number of RPCs), **compound procedures** is applied (in NFS v4), where several RPCs can be grouped into a single request.



Network File System

□ Naming in NFS


- The fundamental idea underlying the NFS naming model is to provide clients **complete transparent access to a remote file system as maintained by a server**. This transparency is achieved by letting a client be able to mount (part of) a remote file system into its own local file system.
- A server exports a directory to a client, and the client mounts the exported directory to its local name space.



Network File System

□ Synchronization in NFS

- Different clients mount the same part of file system in the server → files are shared among the clients.
- How to achieve synchronization if many clients try to read/write the same file?
- If all the reads and writes go directly to the file server, the server can easily achieve the synchronization by **processing them strictly sequentially**.



- However, enabling the server to process every read and write may degrade the performance (in terms of throughput) of NFS, involving too many RPCs.

Network File System

□ Synchronization in NFS

- In order to improve the performance, clients are allowed to maintain local copies of heavily-used files in their private local caches. Thus, clients can read/write their local copies, and flush back the cached data to server when files are closed.
- If a client locally modifies a cached file and shortly thereafter another client reads the file from the server, the second client will get an **obsolete** file.
- There is no way to guarantee clients can read the up-to-date file in NFS.
 - ✓ Rule: changes to an open file are initially visible only to the client that modified the file. Only when the file is closed are the changes made visible to other clients.

